

Test report

1. Database schema:

id	name	continent	area	population	gdp	headOfState	independence	independenceYear	capital	language	official	currency	countryCode
1	1807 West Bank	Asia	584000	5.0M	11.5B	1	1953	1	Jerusalem	Hebrew	Hebrew	9.472B	9472
2	2015 Occupied Golan Heights	Asia	90000	0.5M	2	24	1	1967	1	Spanish	Spanish	34.4 ABG	344
3	2015 Occupied East Jerusalem	Asia	10000	0.5M	3	24	1	1967	1	Hebrew	Hebrew	3.4 ABG	344
4	2015 EastJerusalem West Bank	Asia	1137500	5.0M	4	100	1	1967	1	Arabic	Arabic	65.4 BIL	654
5	2027 Hispania	Europe	744000	ME2	5	600	0	1	1	Latin	Latin	5.4 MBT	54
6	21401 Spain	Europe	505000	45.0M	6	100	1	1978	1	Spanish	Spanish	5.4 BIL	54
7	1044 Iberian	Europe	140000	10.0M	7	970	0	1	1	Spanish	Spanish	4.1 BIL	200
8	474 Asturias	Oceania	200400	8.0M	8	190	0	1	1	Spanish	Spanish	10.3 CSD	103
9	2014 NewCaledonia	Oceania	172400	0.2M	9	120	0	1	1	French	French	7.4 MBT	74
10	2019 Frederiksberg	Europe	60127	1.0M	10	120	1	1	1	Danish	Danish	86.87B	878

2. Queries

- **Most popular languages by region (version which works on every database engine)**

```
SELECT language.name AS language_name, MAX(language_aggregate.region) AS region,
MAX(language_aggregate.language_population) AS language_population FROM (
```

```
SELECT language_sum.region, MAX(language_sum.language_population) AS
language_population FROM (
```

```
SELECT lang.name, cntr.region, sum((lang.percentage / 100) * cntr.population) AS  
language_population
```

FROM language lang

JOIN country cntr

ON lang.country = cntr.code

GROUP BY lang.name, cntr.region) language_sum

GROUP BY language_sum.region) language_aggregate

JOIN country

ON country.region = language_aggregate.region

JOIN language

ON language.country = country.code

GROUP BY language.name, country.region

```
HAVING SUM((language.percentage / 100) * country.population) =  
MAX(language_aggregate.language_population);
```

	LANGUAGE_NAME	REGION	LANGUAGE_POPULATION
1	Javanese	Southeast Asia	83570158
2	Papuan Languages	Melanesia	3792451
3	English	Australia and New Zealand	18695372
4	Russian	Eastern Europe	148136548
5	English	Micronesia/Caribbean	94
6	Arabic	Middle East	89153760
7	German	Western Europe	87155999,6
8	Portuguese	South America	166037997
9	Spanish	Caribbean	21530888
10	English	North America	258821522
11	Chinese	Eastern Asia	1175677671
12	Arabic	Northern Africa	140084044
13	Oromo	Eastern Africa	19395150
14	Kiribati	Micronesia	84235
15	Kongo	Central Africa	11480181
16	Samoan	Polynesia	147108
17	Zulu	Southern Africa	9508689
18	Hausa	Western Africa	29225396
19	Spanish	Central America	122159129
20	Swedish	Nordic Countries	8255142,6
21	English	British Islands	61728266,6
22	Hindi	Southern and Central Asia	405169038
23	Italian	Southern Europe	54303880
24	Lithuanian	Baltic Countries	3047066,4

Database Statistics

POSTGRESQL BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=70, operationsPerSecond=42.857142857142854)
 ORACLE BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=82, operationsPerSecond=36.58536585365854)
 MYSQL BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=151, operationsPerSecond=19.867549668874172)
 H2 BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=1658, operationsPerSecond=1.8094089264173705)

- **Most popular languages by region (version which works only on Oracle and PostgreSQL)**

```
SELECT t.name, t.region, t.language_population
FROM (
  SELECT l.name, c.region, ((l.percentage * c.population)/100) language_population,
  ROW_NUMBER() OVER (PARTITION BY c.region ORDER BY ((l.percentage *
  c.population)/100) DESC) row_number
FROM language l
JOIN country c ON c.code = l.country) t WHERE t.row_number = 1
```

Database Statistics

POSTGRESQL BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=69, operationsPerSecond=43.47826086956521, exception=null)

ORACLE BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=50, operationsPerSecond=60.0, exception=null)

H2 BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=3, operationsPerSecond=1.7976931348623157E308, exception=java.lang.RuntimeException: org.h2.jdbc.JdbcSQLException: Syntax error in SQL statement "SELECT T.NAME, T.REGION, T.LANGUAGE_POPULATION FROM (SELECT L.NAME, C.REGION, ((L.PERCENTAGE * C.POPULATION)/100) LANGUAGE_POPULATION, ROW_NUMBER() OVER (PARTITION[*] BY C.REGION ORDER BY ((L.PERCENTAGE * C.POPULATION)/100) DESC) ROW_NUMBER FROM LANGUAGE L JOIN COUNTRY C ON C.CODE = L.COUNTRY) T WHERE T.ROW_NUMBER = 1 "; expected "; SQL statement: SELECT t.name, t.region, t.language_population FROM (SELECT l.name, c.region, ((l.percentage * c.population)/100) language_population, ROW_NUMBER() OVER (PARTITION BY c.region ORDER BY ((l.percentage * c.population)/100) DESC) row_number FROM language l JOIN country c ON c.code = l.country) t WHERE t.row_number = 1 (42001-197))

MYSQL BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=3, operationsPerSecond=1.7976931348623157E308, exception=java.lang.RuntimeException: com.mysql.jdbc.exceptions.jdbc4.MySQLSyntaxErrorException: You have an error in your SQL syntax; check the manual that corresponds to your MySQL server version for the right syntax to use near 'row_number FROM language l JOIN country c ON c.code = l.country) t WHERE t.ro' at line 3)

- **Most popular languages by region (version for Mongo)**

```
db.language.aggregate([{$lookup:
{from:"country",localField:"country",foreignField:"code",as:"countryObj"}},{$project:{countryObj:
1, name: 1, percentage: 1,totalPopulation: {$multiply: [{ $divide: ["$percentage", 100]}, {
$arrayElemAt: ["$countryObj.population",
1]}]}},{$sort:{"countryObj.region":1,"totalPopulation":1}},{$group:{_id:"countryObj.region",docs:
{ $push:"$$ROOT"}},{$project:{"top_one": {$slice:["$docs", 1]}}}]);
```

Database Statistics

MONGO BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=900, operationsPerSecond=3.333333333333333, exception=null)

- **Cities from 10 biggest countries**

```
SELECT town.*
FROM country
JOIN town
ON town.country = country.code
WHERE (
SELECT COUNT(*)
FROM country cntr
WHERE cntr.surfacearea >= country.surfacearea
) <= 10
ORDER BY town.population DESC;
```

	ID	DISTRICT	NAME	POPULATION	COUNTRY
1	1024	Maharashtra	Mumbai (Bombay)	10500000	IND
2	206	São Paulo	São Paulo	9968485	BRA
3	1890	Shanghai	Shanghai	9696300	CHN
4	3580	Moscow (City)	Moscow	8389200	RUS
5	3793	New York	New York	8008278	USA
6	1891	Peking	Peking	7472000	CHN
7	1025	Delhi	Delhi	7206704	IND
8	1892	Chongqing	Chongqing	6351600	CHN
9	207	Rio de Janeiro	Rio de Janeiro	5598953	BRA
10	1893	Tianjin	Tianjin	5286800	CHN
11	3581	Pietari	St Petersburg	4694000	RUS
12	1026	West Bengali	Calcutta [Kolkata]	4399819	IND
13	1894	Hubei	Wuhan	4344600	CHN
14	1895	Heilongjiang	Harbin	4289800	CHN
15	1896	Liaoning	Shenyang	4265200	CHN
16	1897	Guangdong	Kanton [Guangzhou]	4256300	CHN
17	1027	Tamil Nadu	Chennai (Madras)	3841396	IND
18	3794	California	Los Angeles	3694820	USA
19	1898	Sichuan	Chengdu	3361500	CHN
20	130	New South Wales	Sydney	3276207	AUS
21	69	Distrito Federal	Buenos Aires	2982146	ARG
22	1028	Andhra Pradesh	Hyderabad	2964638	IND
23	3795	Illinois	Chicago	2896016	USA
24	1029	Gujarat	Ahmedabad	2876710	IND
25	1899	Jiangsu	Nanking [Nanjing]	2870300	CHN
26	131	Victoria	Melbourne	2865329	AUS
27	1900	Jilin	Changchun	2812000	CHN
28	1901	Shaanxi	Xi'an	2761400	CHN
29	1902	Liaoning	Dalian	2697000	CHN
30	1030	Karnataka	Bangalore	2660088	IND
31	1903	Shandong	Qingdao	2596000	CHN
32	208	Bahia	Salvador	2302832	BRA
33	1904	Shandong	Jinan	2278100	CHN
34	1905	Zhejiang	Hangzhou	2190500	CHN

Database Statistics
 MYSQL BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=255, operationsPerSecond=11.76470588235294)
 POSTGRESQL BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=325, operationsPerSecond=9.23076923076923)
 H2 BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=1999, operationsPerSecond=1.5007503751875937)
 ORACLE BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=115, operationsPerSecond=26.08695652173913)

- **Cities from 10 biggest countries (mongodb version)**

```
db.country.aggregate([ {$lookup: {from: "town", localField: "code", foreignField: "country", as: "countryTown"}}, {$sort: {surfacearea: -1}}, {$limit: 10} ]]);
```

Database Statistics
 MONGO BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=4483, operationsPerSecond=0.6691947356680794, exception=null)

- **Top 3 cities from every region based on population**

```
SELECT town.*, country.region  
  
FROM town  
  
JOIN country  
  
ON town.country = country.code  
  
WHERE (
```

```

SELECT COUNT(*)

FROM town twn

JOIN country cntr

ON twn.country = cntr.code

WHERE twn.population >= town.population AND country.region = cntr.region

) <= 3

ORDER BY country.region;

```

ID	DISTRICT	NAME	POPULATION	COUNTRY	REGION
1	132 Queensland	Brisbane	1291117	AUS	Australia and New Zealand
2	130 New South Wales	Sydney	3276207	AUS	Australia and New Zealand
3	131 Victoria	Melbourne	2865329	AUS	Australia and New Zealand
4	2434 Riika	Riga	764328	LVA	Baltic Countries
5	2447 Vilna	Vilnius	577969	LTU	Baltic Countries
6	2448 Kaunas	Kaunas	412639	LTU	Baltic Countries
7	457 England	Birmingham	1013000	GBR	British Islands
8	456 England	London	7285000	GBR	British Islands
9	458 Scotland	Glasgow	619680	GBR	British Islands
10	929 Ouest	Port-au-Prince	884472	HTI	Caribbean
11	587 Distrito Nacional	Santo Domingo de Guzmán	1609966	DOM	Caribbean
12	2413 La Habana	La Habana	2256000	CUB	Caribbean
13	56 Luanda	Luanda	2022000	AGO	Central Africa
14	1803 Littoral	Douala	1448300	CMR	Central Africa
15	2298 Kinshasa	Kinshasa	5064000	COD	Central Africa
16	2517 México	Ecatepec de Morelos	1620303	MEX	Central America
17	2515 Distrito Federal	Ciudad de México	8591309	MEX	Central America
18	2516 Jalisco	Guadalajara	1647720	MEX	Central America
19	756 Addis Abeba	Addis Abeba	2495000	ETH	Eastern Africa
20	1881 Nairobi	Nairobi	2290000	KEN	Eastern Africa
21	3305 Dar es Salaam	Dar es Salaam	1747000	TZA	Eastern Africa
22	1532 Tokyo-to	Tokyo	7980230	JPN	Eastern Asia
23	1890 Shanghai	Shanghai	9696300	CHN	Eastern Asia
24	2331 Seoul	Seoul	9981619	KOR	Eastern Asia
25	3580 Moscow (City)	Moscow	8389200	RUS	Eastern Europe
26	3581 Pietari	St Petersburg	4694000	RUS	Eastern Europe
27	3426 Kiöva	Kyiv	2624000	UKR	Eastern Europe
28	3493 -	Nouméa	76293	NCL	Melanesia
29	2884 National Capital Dis	Port Moresby	247000	PNG	Melanesia
30	764 Central	Suva	77366	FJI	Melanesia
31	2881 Koror	Koror	12000	PLW	Micronesia
32	2688 Chuuk	Weno	22000	FSM	Micronesia
33	2507 Majuro	Dalap-Uliga-Darrit	28000	MHL	Micronesia
34	1365 Baghdad	Baghdad	4336000	IRQ	Middle East
35	3173 Riyadh	Riyadh	3324000	SAU	Middle East
36	3357 Istanbul	Istanbul	8787958	TUR	Middle East
37	3236 Newmaa	Helsinki [Helsingfors]	555474	FIN	Nordic Countries

Database Statistics

MYSQL BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=13134, operationsPerSecond=0.2284148012791229)

POSTGRESQL BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=17164, operationsPerSecond=0.17478443253320902)

H2 BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=18952, operationsPerSecond=0.15829463908822286)

ORACLE BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=18904, operationsPerSecond=0.15869657215404148)

- **Top 3 cities from every region based on population (mongo version)**

```

db.town.aggregate([{$lookup:{from:"country",localField:"country",foreignField:"code",as:"countryObj"}},{ $sort:{"countryObj.region":1,"population":1}},{ $group:{_id:"countryObj.region",docs:{$push:"$$ROOT"}},{ $project:{"top_three": {$slice:["$docs", 3]}}}]

```

Database Statistics

MONGO BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=23077, operationsPerSecond=0.9749756256093598, exception=null)

3. Conclusions

Three of the testes databases can be used in production environment, Oracle, MySQL, PostgreSQL. H2 is here only for contrast. As you can see H2 have worst execution times for every query, in first case H2 was about 20 times worst than other databases.

Let's then focus on these three databases mentioned before, in each case different database won, but in general, Oracle had the most stable results, it of course depends on many factors, for example all of these engines works simultaneously on my local machine, which of course is not the best testing environment.

Another case are indexes, in each queries there was some subqueries, this means at least square complexity, indexes on COUNTRY.POPULATION, COUNTRY.SURFACEAREA and TOWN.POPULATION would for sure speed up a bit all the queries, there can be more indexes of course, but I think, these on mentioned numeric values, are most valuable, as I often searched using these values. Additionally, it is difficult to find two same values for mentioned columns, and that's good from index point of view.

Times after indexes was added on mentioned columns:

- **Most popular languages by region (version which works on every database engine)**

Database	Statistics
H2	BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=442, operationsPerSecond=6.787330316742081, exception=null)
ORACLE	BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=50, operationsPerSecond=61.2244897959183672, exception=null)
POSTGRESQL	BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=49, operationsPerSecond=61.224489795918366, exception=null)
MYSQL	BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=80, operationsPerSecond=37.5, exception=null)

- **Most popular languages by region (version which works only on Oracle and PostgreSQL)**

Database	Statistics
H2	BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=1797691148621157E308, exception=java.lang.RuntimeException: org.h2.jdbc.JdbcSQLException: Syntax error in SQL statement "SELECT T.NAME, T.REGION, T.LANGUAGE, POPULATION FROM (SELECT L.NAME, C.REGION, ((L.PERCENTAGE * C.POPULATION)/100) LANGUAGE_POPULATION, ROW_NUMBER() OVER (PARTITION BY C.REGION ORDER BY ((L.PERCENTAGE * C.POPULATION)/100) DESC) ROW_NUMBER FROM LANGUAGE L JOIN COUNTRY C ON C.CODE = L.COUNTRY) T WHERE T.ROW_NUMBER = 1 ", expected 77; SQL statement: SELECT t.name, t.region, t.language, population FROM (SELECT t.name, c.region, ((t.percentage * c.population)/100) language_population, ROW_NUMBER() OVER (PARTITION BY c.region ORDER BY ((t.percentage * c.population)/100) DESC) row_number FROM language l JOIN country c ON c.code = l.country) t WHERE row_number = 1 [42001-1977])
ORACLE	BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=25, operationsPerSecond=120.0, exception=null)
POSTGRESQL	BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=56, operationsPerSecond=53.57142857142857, exception=null)
MYSQL	BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=3, operationsPerSecond=1.797691148621157E308, exception=java.lang.RuntimeException: com.mysql.jdbc.exceptions.jdbc4.MySQLException: You have an error in your SQL syntax; check the manual that corresponds to your MySQL server version for the right syntax to use near 'row_number FROM language l JOIN country c ON c.code = l.country) t WHERE row' at line 3)

- **Cities from 10 biggest countries**

Database	Statistics
H2	BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=150, operationsPerSecond=20.0, exception=null)
ORACLE	BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=65, operationsPerSecond=46.15384615384615, exception=null)
POSTGRESQL	BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=142, operationsPerSecond=21.126760563380284, exception=null)
MYSQL	BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=135, operationsPerSecond=22.22222222222222, exception=null)

- **Top 3 cities from every region based on population**

Database	Statistics
H2	BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=16738, operationsPerSecond=0.1792328832596487, exception=null)
ORACLE	BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=37243, operationsPerSecond=0.08055205004967377, exception=null)
POSTGRESQL	BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=11849, operationsPerSecond=0.2531859228626888, exception=null)
MYSQL	BenchmarkController.BenchmarkStatistics(totalOperations=3, totalOperationsTime=10465, operationsPerSecond=0.2866698518872432, exception=null)

With indexes on tables I had an average 25% time improvement. But of course this have some impact on tables initialization, on every application bootstrap, databases are filled with new data, now it takes more time than before.